

# BIA 5302\_Group Project

August 6, 2023

## 1 Machine Learning and Programming 1 - BIA-5302 - Group Project

### 1.1 Group Members :

1.1.1 Damandeep Singh - N01533482

1.1.2 Ishan Sahni - N01530356

1.1.3 Chandeeep Singh - N01530793

1.1.4 Manoj Maddukuri - N01530629

## 2 MILESTONE 2

### 2.1 Importing libraries, loading and viewing the data

```
[1]: import pandas as pd
import numpy as np
import warnings
warnings.filterwarnings('ignore')
import matplotlib.pyplot as plt
```

```
[2]: cars = pd.read_csv("E:/Documents/Humber/Semester 3/Machine Learning and
↳ Programming 1/Group Project/BIA 5302_Group Project_Data.csv")
```

```
[3]: cars.head()
```

```
[3]:
```

	ID	Price	Levy	Manufacturer	Model	Prod. year	Category	\
0	45654403	13328	1399	LEXUS	RX 450	2010	Jeep	
1	44731507	16621	1018	CHEVROLET	Equinox	2011	Jeep	
2	45774419	8467	-	HONDA	FIT	2006	Hatchback	
3	45769185	3607	862	FORD	Escape	2011	Jeep	
4	45809263	11726	446	HONDA	FIT	2014	Hatchback	

	Leather interior	Fuel type	Engine volume	Mileage	Cylinders	\
0	Yes	Hybrid	3.5	186005 km	6	
1	No	Petrol	3	192000 km	6	
2	No	Petrol	1.3	200000 km	4	
3	Yes	Hybrid	2.5	168966 km	4	

4	Yes	Petrol	1.3	91901 km	4
---	-----	--------	-----	----------	---

	Gear box type	Drive wheels	Doors	Wheel	Color	Airbags
0	Automatic	4x4	04-May	Left wheel	Silver	12
1	Tiptronic	4x4	04-May	Left wheel	Black	8
2	Variator	Front	04-May	Right-hand drive	Black	2
3	Automatic	4x4	04-May	Left wheel	White	0
4	Automatic	Front	04-May	Left wheel	Silver	4

```
[4]: #Checking the datatypes of the variables
cars.dtypes
```

```
[4]: ID                int64
Price                int64
Levy                object
Manufacturer        object
Model              object
Prod. year          int64
Category            object
Leather interior    object
Fuel type           object
Engine volume       object
Mileage             object
Cylinders           int64
Gear box type       object
Drive wheels        object
Doors               object
Wheel              object
Color              object
Airbags            int64
dtype: object
```

```
[5]: #Printing the summary statistics
cars.describe()
```

```
[5]:
```

	ID	Price	Prod. year	Cylinders	Airbags
count	1.923700e+04	1.923700e+04	19237.000000	19237.000000	19237.000000
mean	4.557654e+07	1.855593e+04	2010.912824	4.582991	6.582627
std	9.365914e+05	1.905813e+05	5.668673	1.199933	4.320168
min	2.074688e+07	1.000000e+00	1939.000000	1.000000	0.000000
25%	4.569837e+07	5.331000e+03	2009.000000	4.000000	4.000000
50%	4.577231e+07	1.317200e+04	2012.000000	4.000000	6.000000
75%	4.580204e+07	2.207500e+04	2015.000000	4.000000	12.000000
max	4.581665e+07	2.630750e+07	2020.000000	16.000000	16.000000

## 3 Data Preparation and Cleaning

### 3.0.1 1. Missing Values

```
[6]: cars.isnull().sum()
```

```
[6]: ID                0
     Price             0
     Levy              0
     Manufacturer      0
     Model             0
     Prod. year        0
     Category          0
     Leather interior   0
     Fuel type         0
     Engine volume     0
     Mileage           0
     Cylinders         0
     Gear box type     0
     Drive wheels      0
     Doors             0
     Wheel             0
     Color             0
     Airbags           0
     dtype: int64
```

There are no null/missing values in the dataset.

### 3.0.2 2. Duplicate Data

```
[7]: duplicates = cars.duplicated()
     duplicates_num = duplicates.sum()
     print("Number of duplicate rows in the dataset are", duplicates_num)
```

Number of duplicate rows in the dataset are 313

```
[8]: cars.drop_duplicates(inplace=True)
     cars
```

```
[8]:
```

	ID	Price	Levy	Manufacturer	Model	Prod. year	Category \
0	45654403	13328	1399	LEXUS	RX 450	2010	Jeep
1	44731507	16621	1018	CHEVROLET	Equinox	2011	Jeep
2	45774419	8467	-	HONDA	FIT	2006	Hatchback
3	45769185	3607	862	FORD	Escape	2011	Jeep
4	45809263	11726	446	HONDA	FIT	2014	Hatchback
...	...	...	...	...	...	...	...
19232	45798355	8467	-	MERCEDES-BENZ	CLK 200	1999	Coupe
19233	45778856	15681	831	HYUNDAI	Sonata	2011	Sedan

19234	45804997	26108	836	HYUNDAI	Tucson	2010	Jeep
19235	45793526	5331	1288	CHEVROLET	Captiva	2007	Jeep
19236	45813273	470	753	HYUNDAI	Sonata	2012	Sedan

	Leather interior	Fuel type	Engine volume	Mileage	Cylinders	\
0	Yes	Hybrid	3.5	186005 km	6	
1	No	Petrol	3	192000 km	6	
2	No	Petrol	1.3	200000 km	4	
3	Yes	Hybrid	2.5	168966 km	4	
4	Yes	Petrol	1.3	91901 km	4	
...	...	...	...	...	...	
19232	Yes	CNG	2.0 Turbo	300000 km	4	
19233	Yes	Petrol	2.4	161600 km	4	
19234	Yes	Diesel	2	116365 km	4	
19235	Yes	Diesel	2	51258 km	4	
19236	Yes	Hybrid	2.4	186923 km	4	

	Gear box type	Drive wheels	Doors	Wheel	Color	Airbags
0	Automatic	4x4	04-May	Left wheel	Silver	12
1	Tiptronic	4x4	04-May	Left wheel	Black	8
2	Variator	Front	04-May	Right-hand drive	Black	2
3	Automatic	4x4	04-May	Left wheel	White	0
4	Automatic	Front	04-May	Left wheel	Silver	4
...	...	...	...	...	...	
19232	Manual	Rear	02-Mar	Left wheel	Silver	5
19233	Tiptronic	Front	04-May	Left wheel	Red	8
19234	Automatic	Front	04-May	Left wheel	Grey	4
19235	Automatic	Front	04-May	Left wheel	Black	4
19236	Automatic	Front	04-May	Left wheel	White	12

[18924 rows x 18 columns]

We have dropped 313 rows from the dataset as these rows would not add anything to our learning process of the model.

### 3.0.3 3. Irrelevant and Incorrect Data

```
[9]: #Displaying the number of "-" values in the 'Levy', i.e., Tax variable
cars['Levy'].describe()
```

```
[9]: count      18924
unique         559
top            -
freq          5709
Name: Levy, dtype: object
```

```
[10]: #Replacing the "-" values to be blank or NULL  
cars['Levy'].replace({'-':np.nan}, inplace=True)
```

```
[11]: cars.isnull().sum()
```

```
[11]: ID                0  
Price                0  
Levy                5709  
Manufacturer         0  
Model                0  
Prod. year           0  
Category             0  
Leather interior     0  
Fuel type            0  
Engine volume        0  
Mileage              0  
Cylinders            0  
Gear box type        0  
Drive wheels         0  
Doors                0  
Wheel               0  
Color               0  
Airbags             0  
dtype: int64
```

```
[12]: #Converting the datatype of 'Levy' from object to float  
cars['Levy'] = cars['Levy'].astype(float)
```

```
[13]: # Replacing the NULL values in 'Levy' with the mean  
mean_levy = cars['Levy'].mean()  
mean_levy
```

```
[13]: 906.2992054483541
```

```
[14]: cars['Levy'].fillna(mean_levy, inplace=True)
```

The incorrect data in 'Levy' variable has been replaced with its mean.

```
[15]: #Changing the 'Doors' format into appropriate numbers  
cars['Doors'].value_counts()  
cars['Doors'].replace({'04-May':4, '02-Mar':2, '>5':5}, inplace=True)
```

The 'Doors' variable has now been formatted in the correct way.

```
[16]: #Removing the "km" unit from the 'Mileage' variable  
cars['Mileage']=cars['Mileage'].str.replace('km', '')
```

```
[17]: cars['Mileage'].value_counts()
```

```
[17]: 0          714
      200000    181
      150000    159
      160000    120
      180000    117
      ...
      100563     1
      354300     1
      21178      1
      110539     1
      186923     1
      Name: Mileage, Length: 7687, dtype: int64
```

```
[18]: #Converting the datatype of 'Mileage'
      cars['Mileage']=cars['Mileage'].astype(int)
```

```
[19]: cars['Mileage'].head()
```

```
[19]: 0    186005
      1    192000
      2    200000
      3    168966
      4     91901
      Name: Mileage, dtype: int32
```

```
[20]: cars['Mileage'].mean()
```

```
[20]: 1555372.718928345
```

The 'Mileage' variable has now been formatted as per the requirements of our model.

```
[21]: cars['Engine volume'].value_counts()
```

```
[21]: 2          3856
      2.5        2246
      1.8        1743
      1.6        1446
      1.5        1289
      ...
      6.8          1
      6.7          1
      3.1          1
      0.8 Turbo     1
      1.1 Turbo     1
      Name: Engine volume, Length: 107, dtype: int64
```

```
[22]: #Removing "Turbo" from the values
      cars['Engine volume']=cars['Engine volume'].str.split(' ').str.get(0)
```

```
[23]: #Converting the datatype of 'Engine Volume'
cars['Engine volume']=cars['Engine volume'].astype(float)
```

```
[24]: cars['Engine volume'].mean()
```

```
[24]: 2.306251321073769
```

The 'Engine volume' variable has now been formatted as per the requirements of our model.

```
[25]: cars['Wheel'].value_counts()
```

```
[25]: Left wheel          17471
Right-hand drive      1453
Name: Wheel, dtype: int64
```

```
[26]: cars['Wheel'].replace({'Right-hand drive':'Right wheel'}, inplace=True)
```

The 'Wheel' variable has now been formatted correctly.

```
[27]: #Dropping the irrelevant record in Manufacturer = áf;áf@áf•áf using boolean
      ↳ indexing
cars = cars[cars['ID']!=45779593]
cars = cars[cars['ID']!=39223518]
```

```
[28]: cars
```

```
[28]:
```

	ID	Price	Levy	Manufacturer	Model	Prod. year	\
0	45654403	13328	1399.000000	LEXUS	RX 450	2010	
1	44731507	16621	1018.000000	CHEVROLET	Equinox	2011	
2	45774419	8467	906.299205	HONDA	FIT	2006	
3	45769185	3607	862.000000	FORD	Escape	2011	
4	45809263	11726	446.000000	HONDA	FIT	2014	
...	...	...	...	...	...	...	
19232	45798355	8467	906.299205	MERCEDES-BENZ	CLK 200	1999	
19233	45778856	15681	831.000000	HYUNDAI	Sonata	2011	
19234	45804997	26108	836.000000	HYUNDAI	Tucson	2010	
19235	45793526	5331	1288.000000	CHEVROLET	Captiva	2007	
19236	45813273	470	753.000000	HYUNDAI	Sonata	2012	

	Category	Leather interior	Fuel type	Engine volume	Mileage	\
0	Jeep	Yes	Hybrid	3.5	186005	
1	Jeep	No	Petrol	3.0	192000	
2	Hatchback	No	Petrol	1.3	200000	
3	Jeep	Yes	Hybrid	2.5	168966	
4	Hatchback	Yes	Petrol	1.3	91901	
...	...	...	...	...	...	
19232	Coupe	Yes	CNG	2.0	300000	
19233	Sedan	Yes	Petrol	2.4	161600	

19234	Jeep	Yes	Diesel	2.0	116365
19235	Jeep	Yes	Diesel	2.0	51258
19236	Sedan	Yes	Hybrid	2.4	186923

	Cylinders	Gear box type	Drive wheels	Doors	Wheel	Color	\
0	6	Automatic	4x4	4	Left wheel	Silver	
1	6	Tiptronic	4x4	4	Left wheel	Black	
2	4	Variator	Front	4	Right wheel	Black	
3	4	Automatic	4x4	4	Left wheel	White	
4	4	Automatic	Front	4	Left wheel	Silver	
...	...	...	...	...	...	...	
19232	4	Manual	Rear	2	Left wheel	Silver	
19233	4	Tiptronic	Front	4	Left wheel	Red	
19234	4	Automatic	Front	4	Left wheel	Grey	
19235	4	Automatic	Front	4	Left wheel	Black	
19236	4	Automatic	Front	4	Left wheel	White	

	Airbags
0	12
1	8
2	2
3	0
4	4
...	...
19232	5
19233	8
19234	4
19235	4
19236	12

[18922 rows x 18 columns]

```
[29]: #Saving the cleaned dataset (so far) in a new dataframe variable
cars_2 = cars.copy()
```

### 3.0.4 4. Categorical Data

```
[30]: #Categorising the data
cars['Wheel'].replace({'Left wheel':1,'Right wheel':0}, inplace=True)
cars['Leather interior'].replace({'Yes':1,'No':0}, inplace=True)
```

```
[31]: cars.head()
```

```
[31]:      ID  Price      Levy Manufacturer  Model  Prod. year  Category \
0  45654403  13328  1399.000000      LEXUS  RX 450      2010      Jeep
1  44731507  16621  1018.000000  CHEVROLET  Equinox      2011      Jeep
2  45774419   8467   906.299205      HONDA    FIT      2006  Hatchback
```



3	45769185	3607	862.000000	FORD	Escape	2011	Jeep
4	45809263	11726	446.000000	HONDA	FIT	2014	Hatchback

	Leather interior	Fuel type	Engine volume	Mileage	Cylinders	\
0	1	Hybrid	3.5	186005	6	
1	0	Petrol	3.0	192000	6	
2	0	Petrol	1.3	200000	4	
3	1	Hybrid	2.5	168966	4	
4	1	Petrol	1.3	91901	4	

	Gear box type	Drive wheels	Doors	Wheel	Color	Airbags
0	Automatic	4x4	4	1	Silver	12
1	Tiptronic	4x4	4	1	Black	8
2	Variator	Front	4	0	Black	2
3	Automatic	4x4	4	1	White	0
4	Automatic	Front	4	1	Silver	4

```
[32]: #Categorising more data
cars['Fuel type'].replace({'Plug-in Hybrid':6,'Petrol':5,'LPG':4,'Hydrogen':
    ↳3,'Hybrid':2,'Diesel':1,'CNG':0}, inplace=True)
cars['Gear box type'].replace({'Variator':3,'Tiptronic':2,'Manual':
    ↳1,'Automatic':0}, inplace=True)
cars['Drive wheels'].replace({'Rear':2,'Front':1,'4x4':0}, inplace=True)
cars['Category'].replace({'Universal':10,'Sedan':9,'Pickup':8,'Minivan':
    ↳7,'Microbus':6,'Limousine':5,'Jeep':4,'Hatchback':3,'Goods wagon':2,'Coupe':
    ↳1,'Cabriolet':0}, inplace=True)
```

```
[33]: #Changing the datatypes for the converted categorical variables
cars['Leather interior'] = cars['Leather interior'].astype(int)
cars['Wheel'] = cars['Wheel'].astype(int)
cars['Fuel type'] = cars['Fuel type'].astype(int)
cars['Gear box type'] = cars['Gear box type'].astype(int)
cars['Drive wheels'] = cars['Drive wheels'].astype(int)
cars['Category'] = cars['Category'].astype(int)
```

```
[34]: cars.dtypes
```

```
[34]: ID                int64
Price                int64
Levy                float64
Manufacturer         object
Model                object
Prod. year           int64
Category             int32
Leather interior     int32
Fuel type            int32
Engine volume        float64
```

```

Mileage          int32
Cylinders        int64
Gear box type    int32
Drive wheels     int32
Doors            int64
Wheel            int32
Color            object
Airbags          int64
dtype: object

```

### 3.0.5 5. Outliers

```
[35]: import plotly.express as px
```

```
[36]: #Checking for the outliers in 'Price' using the quartile method
q1, q3 = np.percentile(cars["Price"], [25, 75])
iqr = q3 - q1
lower_bound = q1 - 1.5*iqr
upper_bound = q3 + 1.5*iqr

#Creating conditions to isolate the outliers
outliers_price_1 = cars["Price"][(cars["Price"] < lower_bound) | (cars["Price"] >
    ↳ upper_bound)]
```

```
[37]: outliers_price_1
```

```
[37]: 14      59464
      36      51746
      47      55390
      56      87112
      73      53154
      ...
     19144     56814
     19161     64290
     19180     63886
     19188     61154
     19211     50037
Name: Price, Length: 1055, dtype: int64
```

```
[38]: mean_Price_without_outliers_1 = cars["Price"][(cars["Price"] >= lower_bound) &
    ↳ (cars["Price"] <= upper_bound)].mean()

#Replacing the outliers with the mean_Price_without_outliers
cars["Price"] = cars["Price"].where(~((cars["Price"] < lower_bound) |
    ↳ (cars["Price"] > upper_bound)), mean_Price_without_outliers_1)
```

```
[39]: #Checking for the outliers in 'Mileage' using the quartile method
q1, q3 = np.percentile(cars["Mileage"], [25, 75])
iqr = q3 - q1
lower_bound = q1 - 1.5*iqr
upper_bound = q3 + 1.5*iqr

#Creating conditions to isolate the outliers
outliers_Mileage = cars["Mileage"][(cars["Mileage"] < lower_bound) |
↪(cars["Mileage"] > upper_bound)]
```

```
[40]: mean_without_outliers = cars["Mileage"][(cars["Mileage"] >= lower_bound) &
↪(cars["Mileage"] <= upper_bound)].mean()
mean_without_outliers
```

```
[40]: 129507.99163340077
```

```
[41]: #Replacing the outliers with the mean_without_outliers
cars["Mileage"] = cars["Mileage"].where(~((cars["Mileage"] < lower_bound) |
↪(cars["Mileage"] > upper_bound)), mean_without_outliers)
```

```
[42]: cars["Mileage"].mean()
```

```
[42]: 129507.99163340076
```

```
[43]: cars.head(15)
```

```
[43]:
```

	ID	Price	Levy	Manufacturer	Model	Prod. year	\
0	45654403	13328.000000	1399.000000	LEXUS	RX 450	2010	
1	44731507	16621.000000	1018.000000	CHEVROLET	Equinox	2011	
2	45774419	8467.000000	906.299205	HONDA	FIT	2006	
3	45769185	3607.000000	862.000000	FORD	Escape	2011	
4	45809263	11726.000000	446.000000	HONDA	FIT	2014	
5	45802912	39493.000000	891.000000	HYUNDAI	Santa FE	2016	
6	45656768	1803.000000	761.000000	TOYOTA	Prius	2010	
7	45816158	549.000000	751.000000	HYUNDAI	Sonata	2013	
8	45641395	1098.000000	394.000000	TOYOTA	Camry	2014	
9	45756839	26657.000000	906.299205	LEXUS	RX 350	2007	
10	45621750	941.000000	1053.000000	MERCEDES-BENZ	E 350	2014	
11	45814819	8781.000000	906.299205	FORD	Transit	1999	
12	45815568	3000.000000	906.299205	OPEL	Vectra	1997	
13	45661288	1019.000000	1055.000000	LEXUS	RX 450	2013	
14	45732604	14039.732636	891.000000	HYUNDAI	Santa FE	2016	

	Category	Leather interior	Fuel type	Engine volume	Mileage	\
0	4	1	2	3.5	186005.000000	
1	4	0	5	3.0	192000.000000	
2	3	0	5	1.3	200000.000000	

3	4	1	2	2.5	168966.000000
4	3	1	5	1.3	91901.000000
5	4	1	1	2.0	160931.000000
6	3	1	2	1.8	258909.000000
7	9	1	5	2.4	216118.000000
8	9	1	2	2.5	129507.991633
9	4	1	5	3.5	128500.000000
10	9	1	1	3.5	184467.000000
11	6	0	0	4.0	0.000000
12	2	0	0	1.6	350000.000000
13	4	1	2	3.5	138038.000000
14	4	1	1	2.0	76000.000000

	Cylinders	Gear box type	Drive wheels	Doors	Wheel	Color	Airbags
0	6	0	0	4	1	Silver	12
1	6	2	0	4	1	Black	8
2	4	3	1	4	0	Black	2
3	4	0	0	4	1	White	0
4	4	0	1	4	1	Silver	4
5	4	0	1	4	1	White	4
6	4	0	1	4	1	White	12
7	4	0	1	4	1	Grey	12
8	4	0	1	4	1	Black	12
9	6	0	0	4	1	Silver	12
10	6	0	2	4	1	White	12
11	8	1	2	2	1	Blue	0
12	4	1	1	4	1	White	4
13	6	0	1	4	1	White	12
14	4	0	1	4	1	White	4

The outliers have now been replaced with the mean values without the outliers.

```
[44]: cars.describe()
```

```
[44]:
```

	ID	Price	Levy	Prod. year	Category \
count	1.892200e+04	18922.000000	18922.000000	18922.000000	18922.000000
mean	4.557571e+07	14039.732636	906.299205	2010.914755	6.266938
std	9.364573e+05	11062.279973	387.172475	5.665814	2.792043
min	2.074688e+07	1.000000	87.000000	1939.000000	0.000000
25%	4.569503e+07	5331.000000	730.000000	2009.000000	4.000000
50%	4.577191e+07	13172.000000	906.299205	2012.000000	7.000000
75%	4.580174e+07	19444.000000	917.000000	2015.000000	9.000000
max	4.581665e+07	47120.000000	11714.000000	2020.000000	10.000000

	Leather interior	Fuel type	Engine volume	Mileage \
count	18922.000000	18922.000000	18922.000000	18922.000000
mean	0.725610	3.427016	2.306252	129507.991633
std	0.446218	1.806268	0.877637	80134.138864

min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	2.000000	1.800000	70195.250000
50%	1.000000	5.000000	2.000000	126400.000000
75%	1.000000	5.000000	2.500000	179200.000000
max	1.000000	6.000000	20.000000	367053.000000

	Cylinders	Gear box type	Drive wheels	Doors	Wheel \
count	18922.000000	18922.000000	18922.000000	18922.000000	18922.000000
mean	4.580277	0.537522	0.909576	3.925378	0.923211
std	1.200271	0.897231	0.566505	0.404158	0.266263
min	1.000000	0.000000	0.000000	2.000000	0.000000
25%	4.000000	0.000000	1.000000	4.000000	1.000000
50%	4.000000	0.000000	1.000000	4.000000	1.000000
75%	4.000000	1.000000	1.000000	4.000000	1.000000
max	16.000000	3.000000	2.000000	5.000000	1.000000

Airbags	
count	18922.000000
mean	6.568914
std	4.322234
min	0.000000
25%	4.000000
50%	6.000000
75%	12.000000
max	16.000000

We now have a new mean in the 'Price' and 'Mileage' variables.

### 3.0.6 6. Feature Scaling

```
[45]: #Feature Scaling on numerical variables won't be applicable since we can't have
      ↪ a fraction for the
      #number of car doors/number of airbags/price a car has. Moreover, since our
      ↪ main business problem relates to predicting
      #the car price(s) using linear regression, Feature Scaling is not necessary.
```

```
[46]: #Since the 'Mileage' variable is in km units, we can change it into Miles
cars["Mileage"] = cars["Mileage"] / 1.609
```

```
[47]: cars
```

	ID	Price	Levy	Manufacturer	Model	Prod. year	\
0	45654403	13328.0	1399.000000	LEXUS	RX 450	2010	
1	44731507	16621.0	1018.000000	CHEVROLET	Equinox	2011	
2	45774419	8467.0	906.299205	HONDA	FIT	2006	
3	45769185	3607.0	862.000000	FORD	Escape	2011	
4	45809263	11726.0	446.000000	HONDA	FIT	2014	

...	...	...	...	...	...	...	
19232	45798355	8467.0	906.299205	MERCEDES-BENZ	CLK 200		1999
19233	45778856	15681.0	831.000000	HYUNDAI	Sonata		2011
19234	45804997	26108.0	836.000000	HYUNDAI	Tucson		2010
19235	45793526	5331.0	1288.000000	CHEVROLET	Captiva		2007
19236	45813273	470.0	753.000000	HYUNDAI	Sonata		2012

	Category	Leather interior	Fuel type	Engine volume	Mileage \
0	4	1	2	3.5	115602.858919
1	4	0	5	3.0	119328.775637
2	3	0	5	1.3	124300.807955
3	4	1	2	2.5	105013.051585
4	3	1	5	1.3	57116.842759

...	...	...	...	...	...	...	
19232	1	1	0	2.0	186451.211933		
19233	9	1	5	2.4	100435.052828		
19234	4	1	1	2.0	72321.317589		
19235	4	1	1	2.0	31857.054071		
19236	9	1	2	2.4	116173.399627		

	Cylinders	Gear box type	Drive wheels	Doors	Wheel	Color	Airbags
0	6	0	0	4	1	Silver	12
1	6	2	0	4	1	Black	8
2	4	3	1	4	0	Black	2
3	4	0	0	4	1	White	0
4	4	0	1	4	1	Silver	4

...	...	...	...	...	...	...	
19232	4	1	2	2	1	Silver	5
19233	4	2	1	4	1	Red	8
19234	4	0	1	4	1	Grey	4
19235	4	0	1	4	1	Black	4
19236	4	0	1	4	1	White	12

[18922 rows x 18 columns]

### 3.0.7 7. Feature Engineering & Selection aka EDA

```
[48]: import seaborn as sns
```

```
[49]: cars.columns
```

```
[49]: Index(['ID', 'Price', 'Levy', 'Manufacturer', 'Model', 'Prod. year',
          'Category', 'Leather interior', 'Fuel type', 'Engine volume', 'Mileage',
          'Cylinders', 'Gear box type', 'Drive wheels', 'Doors', 'Wheel', 'Color',
          'Airbags'],
          dtype='object')
```

```
[50]: year_intervals = [1930,1940,1950,1960,1970,1980,1990,2000,2010,2020]

#Creating a new column 'Year Interval' using pd.cut()
cars['Year Interval'] = pd.cut(cars['Prod. year'], bins=year_intervals,
    ↪labels=[1930,1940,1950,1960,1970,1980,1990,2000,2010])
```

```
[51]: cars['Year Interval'] = cars['Year Interval'].astype(int)
```

```
[52]: cars.dtypes
```

```
[52]: ID                int64
Price                float64
Levy                float64
Manufacturer         object
Model               object
Prod. year          int64
Category            int32
Leather interior     int32
Fuel type           int32
Engine volume       float64
Mileage             float64
Cylinders           int64
Gear box type       int32
Drive wheels        int32
Doors               int64
Wheel              int32
Color               object
Airbags            int64
Year Interval       int32
dtype: object
```

```
[53]: #Dropping the columns that we don't need for our analysis
cars = cars.drop(['ID', 'Manufacturer', 'Model', 'Color', 'Prod. year'], axis =
    ↪1)
cars
```

```
[53]:
```

	Price	Levy	Category	Leather interior	Fuel type	\
0	13328.0	1399.000000	4	1	2	
1	16621.0	1018.000000	4	0	5	
2	8467.0	906.299205	3	0	5	
3	3607.0	862.000000	4	1	2	
4	11726.0	446.000000	3	1	5	
...	...	...	...	...	...	
19232	8467.0	906.299205	1	1	0	
19233	15681.0	831.000000	9	1	5	
19234	26108.0	836.000000	4	1	1	
19235	5331.0	1288.000000	4	1	1	

19236	470.0	753.000000	9	1	2
-------	-------	------------	---	---	---

	Engine volume	Mileage	Cylinders	Gear box type	Drive wheels \
0	3.5	115602.858919	6	0	0
1	3.0	119328.775637	6	2	0
2	1.3	124300.807955	4	3	1
3	2.5	105013.051585	4	0	0
4	1.3	57116.842759	4	0	1
...	...	...	...	...	...
19232	2.0	186451.211933	4	1	2
19233	2.4	100435.052828	4	2	1
19234	2.0	72321.317589	4	0	1
19235	2.0	31857.054071	4	0	1
19236	2.4	116173.399627	4	0	1

	Doors	Wheel	Airbags	Year Interval
0	4	1	12	2000
1	4	1	8	2010
2	4	0	2	2000
3	4	1	0	2010
4	4	1	4	2010
...	...	...	...	...
19232	2	1	5	1990
19233	4	1	8	2010
19234	4	1	4	2000
19235	4	1	4	2000
19236	4	1	12	2010

[18922 rows x 14 columns]

```
[54]: #Creating a Correlation Matrix to identify the relationships
cor_matrix = cars.corr()
cor_matrix
```

```
[54]:
```

	Price	Levy	Category	Leather interior	Fuel type \
Price	1.000000	-0.049339	-0.054161	0.071535	-0.077177
Levy	-0.049339	1.000000	-0.048773	0.011110	0.066871
Category	-0.054161	-0.048773	1.000000	0.092222	0.112219
Leather interior	0.071535	0.011110	0.092222	1.000000	-0.033962
Fuel type	-0.077177	0.066871	0.112219	-0.033962	1.000000
Engine volume	0.012611	0.537900	0.003950	0.271770	0.022790
Mileage	-0.149281	0.073371	-0.020399	-0.000446	-0.145110
Cylinders	-0.027497	0.459809	-0.064786	0.199709	0.078432
Gear box type	0.126295	0.003041	-0.009389	-0.288332	0.103600
Drive wheels	0.025958	-0.121991	0.212471	-0.087913	-0.041288
Doors	0.024269	-0.040903	0.237971	0.106839	-0.045759
Wheel	0.138458	-0.116834	0.119425	0.346664	-0.081819




Airbags	-0.061931	0.067882	0.134575	0.161838	0.078001
Year Interval	0.220925	-0.193565	0.060869	0.347812	0.051609

	Engine volume	Mileage	Cylinders	Gear box type \
Price	0.012611	-0.149281	-0.027497	0.126295
Levy	0.537900	0.073371	0.459809	0.003041
Category	0.003950	-0.020399	-0.064786	-0.009389
Leather interior	0.271770	-0.000446	0.199709	-0.288332
Fuel type	0.022790	-0.145110	0.078432	0.103600
Engine volume	1.000000	0.175918	0.777237	-0.008490
Mileage	0.175918	1.000000	0.155393	0.001147
Cylinders	0.777237	0.155393	1.000000	0.060021
Gear box type	-0.008490	0.001147	0.060021	1.000000
Drive wheels	-0.222435	-0.072303	-0.206221	0.088663
Doors	-0.017697	0.000345	-0.036024	-0.087742
Wheel	0.185996	0.006190	0.091971	-0.136051
Airbags	0.222006	-0.022878	0.174067	0.109907
Year Interval	-0.038372	-0.185211	-0.099701	-0.179269

	Drive wheels	Doors	Wheel	Airbags	Year Interval
Price	0.025958	0.024269	0.138458	-0.061931	0.220925
Levy	-0.121991	-0.040903	-0.116834	0.067882	-0.193565
Category	0.212471	0.237971	0.119425	0.134575	0.060869
Leather interior	-0.087913	0.106839	0.346664	0.161838	0.347812
Fuel type	-0.041288	-0.045759	-0.081819	0.078001	0.051609
Engine volume	-0.222435	-0.017697	0.185996	0.222006	-0.038372
Mileage	-0.072303	0.000345	0.006190	-0.022878	-0.185211
Cylinders	-0.206221	-0.036024	0.091971	0.174067	-0.099701
Gear box type	0.088663	-0.087742	-0.136051	0.109907	-0.179269
Drive wheels	1.000000	-0.145813	0.008274	-0.020755	-0.114501
Doors	-0.145813	1.000000	0.013542	0.048115	0.170801
Wheel	0.008274	0.013542	1.000000	0.146663	0.166060
Airbags	-0.020755	0.048115	0.146663	1.000000	0.212378
Year Interval	-0.114501	0.170801	0.166060	0.212378	1.000000

```
[55]: #Plotting a heatmap (multivariate)
import seaborn as sns

fig, ax = plt.subplots(figsize=(20, 12))
sns.heatmap(cor_matrix, annot=True, cmap='coolwarm', ax=ax)
ax.set_title("Correlation Matrix ")
plt.show()
```



We can see that ‘Cylinders’ and ‘Engine Volume’ have a high correlation - This shows that a car with a bigger engine can hold more cylinders in its capacity.

```
[56]: cars.columns
```

```
[56]: Index(['Price', 'Levy', 'Category', 'Leather interior', 'Fuel type',
        'Engine volume', 'Mileage', 'Cylinders', 'Gear box type',
        'Drive wheels', 'Doors', 'Wheel', 'Airbags', 'Year Interval'],
        dtype='object')
```

```
[57]: #Dividing into Predictors and Outcome(s)
#Feature Selection will be based on the demographic information and the
      ↳relationships mentioned in the question
predictors = ['Levy', 'Category', 'Leather interior', 'Engine volume',
      ↳'Mileage', 'Fuel type', 'Gear box type', 'Drive wheels',
        'Cylinders', 'Doors', 'Wheel', 'Airbags', 'Year Interval']
outcome = 'Price'
```

```
[58]: # Partitioning the data into predictors (X) and output (Y)
X = pd.get_dummies(cars[predictors], drop_first=True)
Y = cars[outcome]
```

### 3.0.8 8. Validation Split

```
[59]: from sklearn.model_selection import train_test_split

      from sklearn.linear_model import LinearRegression

      from dmba import regressionSummary, exhaustive_search
      from dmba import backward_elimination, forward_selection, stepwise_selection
      from dmba import adjusted_r2_score, AIC_score, BIC_score

[60]: #Splitting the data into training and test set
      train_X, valid_X, train_Y, valid_Y = train_test_split(X, Y, test_size = 0.4 ,
      ↪random_state = 3)

[61]: #Performing backward elimination feature selection
      def train_model(variables):
          model = LinearRegression()
          model.fit(train_X[variables], train_Y)
          return model

      def score_model(model, variables):
          return AIC_score(train_Y, model.predict(train_X[variables]), model)

      best_model, best_variables = backward_elimination(train_X.columns, train_model,
      ↪score_model, verbose=True)

      print(best_variables)

      #Performing backward elimination feature selection regressionSummary(valid_Y,
      ↪best_model.predict(valid_X[best_variables]))

Variables: Levy, Category, Leather interior, Engine volume, Mileage, Fuel type,
Gear box type, Drive wheels, Cylinders, Doors, Wheel, Airbags, Year Interval
Start: score=241787.05
Step: score=241785.66, remove Levy
Step: score=241784.82, remove Leather interior
Step: score=241784.82, remove None
['Category', 'Engine volume', 'Mileage', 'Fuel type', 'Gear box type', 'Drive
wheels', 'Cylinders', 'Doors', 'Wheel', 'Airbags', 'Year Interval']

[62]: #Performing forward feature selection
      def train_model(variables):
          if len(variables) == 0:
              return None
          model = LinearRegression()
          model.fit(train_X[variables], train_Y)
          return model
```

```

def score_model(model, variables):
    if len(variables) == 0:
        return AIC_score(train_Y, [train_Y.mean()] * len(train_Y), model, df=1)
    return AIC_score(train_Y, model.predict(train_X[variables]), model)

best_model, best_variables = forward_selection(train_X.columns, train_model,
↪score_model, verbose=True)

print(best_variables)

regressionSummary(valid_Y, best_model.predict(valid_X[best_variables]))

```

Variables: Levy, Category, Leather interior, Engine volume, Mileage, Fuel type, Gear box type, Drive wheels, Cylinders, Doors, Wheel, Airbags, Year Interval  
Start: score=243663.84, constant  
Step: score=243081.34, add Year Interval  
Step: score=242730.10, add Gear box type  
Step: score=242497.39, add Airbags  
Step: score=242224.38, add Wheel  
Step: score=242082.13, add Mileage  
Step: score=241917.37, add Fuel type  
Step: score=241875.51, add Category  
Step: score=241840.75, add Engine volume  
Step: score=241811.19, add Drive wheels  
Step: score=241792.19, add Cylinders  
Step: score=241784.82, add Doors  
Step: score=241784.82, add None  
['Year Interval', 'Gear box type', 'Airbags', 'Wheel', 'Mileage', 'Fuel type', 'Category', 'Engine volume', 'Drive wheels', 'Cylinders', 'Doors']

Regression statistics

Mean Error (ME) : -68.2835  
Root Mean Squared Error (RMSE) : 10237.5982  
Mean Absolute Error (MAE) : 7977.0347  
Mean Percentage Error (MPE) : -1583.4091  
Mean Absolute Percentage Error (MAPE) : 1619.2330

Thus, it can be seen that ‘Levy’ and ‘Leather interior’ must be removed from our linear regression model in order to have the algorithm predict car price(s) more accurately.

### 3.0.9 9. Exploratory Data Analysis (EDA)

```
[63]: cars_2.head()
```

```

[63]:      ID  Price      Levy  Manufacturer  Model  Prod. year  Category \
0  45654403  13328  1399.000000      LEXUS    RX 450        2010      Jeep

```

1	44731507	16621	1018.000000	CHEVROLET	Equinox	2011	Jeep
2	45774419	8467	906.299205	HONDA	FIT	2006	Hatchback
3	45769185	3607	862.000000	FORD	Escape	2011	Jeep
4	45809263	11726	446.000000	HONDA	FIT	2014	Hatchback

	Leather interior	Fuel type	Engine volume	Mileage	Cylinders	Gear box type	\
0	Yes	Hybrid	3.5	186005	6	Automatic	
1	No	Petrol	3.0	192000	6	Tiptronic	
2	No	Petrol	1.3	200000	4	Variator	
3	Yes	Hybrid	2.5	168966	4	Automatic	
4	Yes	Petrol	1.3	91901	4	Automatic	

	Drive wheels	Doors	Wheel	Color	Airbags
0	4x4	4	Left wheel	Silver	12
1	4x4	4	Left wheel	Black	8
2	Front	4	Right wheel	Black	2
3	4x4	4	Left wheel	White	0
4	Front	4	Left wheel	Silver	4

```
[64]: #Analyzing the cars sold based on color types (univariate)

#Categorising and clubbing a few car colors into one 'Other'
cars_2['Color'].replace({'Purple' : 'Other', 'Yellow' : 'Other', 'Sky blue' :
↳ 'Other',
                        'Golden' : 'Other', 'Carnelian red' : 'Other', 'Beige' :
↳ 'Other', 'Brown' : 'Other'}, inplace=True)
color_counts = cars_2['Color'].value_counts()
color_counts
```

```
[64]: Black      4944
      White      4406
      Silver     3728
      Grey       2343
      Blue       1376
      Other       905
      Red         622
      Green       321
      Orange      252
      Pink        25
      Name: Color, dtype: int64
```

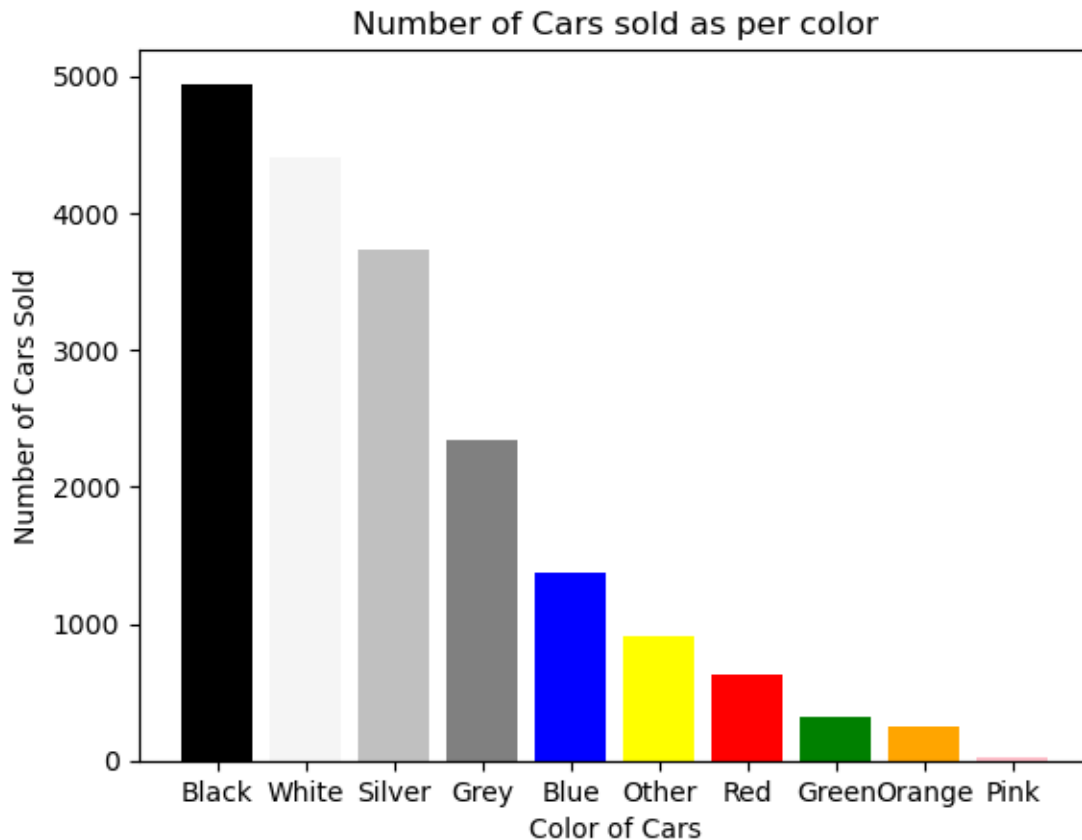
```
[65]: #Creating a custom color palette

color_palette = ['black', 'whitesmoke', 'silver', 'grey', 'blue', 'yellow',
↳ 'red', 'green', 'orange', 'pink']

plt.bar(color_counts.index, color_counts.values, color = color_palette)
```

```
plt.xlabel('Color of Cars')
plt.ylabel('Number of Cars Sold')
plt.title('Number of Cars sold as per color')

plt.show()
```



As seen, “Black” is the most preferred color for car buyers, followed by “White” and “Silver”.

```
[66]: import seaborn as sns
```

```
[67]: #Analyzing the cars sold based on the their Category (univariate)
```


```
#Categorising and clubbing a few car categories into one 'Other'
cars_2['Category'].replace({'Cabriolet' : 'Other', 'Pickup' : 'Other', 'Goods_
↪wagon' : 'Other',
                           'Universal' : 'Other', 'Limousine' : 'Other',
↪'Microbus' : 'Other'}, inplace=True)

Categories = cars_2['Category'].value_counts()
```

Categories

```
[67]: Sedan      8600  
      Jeep      5377  
      Hatchback  2799  
      Other      985  
      Minivan    633  
      Coupe      528  
      Name: Category, dtype: int64
```

```
[68]: palette_color = sns.color_palette('bright')  
      plt.pie(Categories, labels = Categories.index, colors = palette_color, autopct=  
              ↳ '%.0f%%')  
      plt.title('Category Distribution of the cars sold')  
  
      plt.show()
```

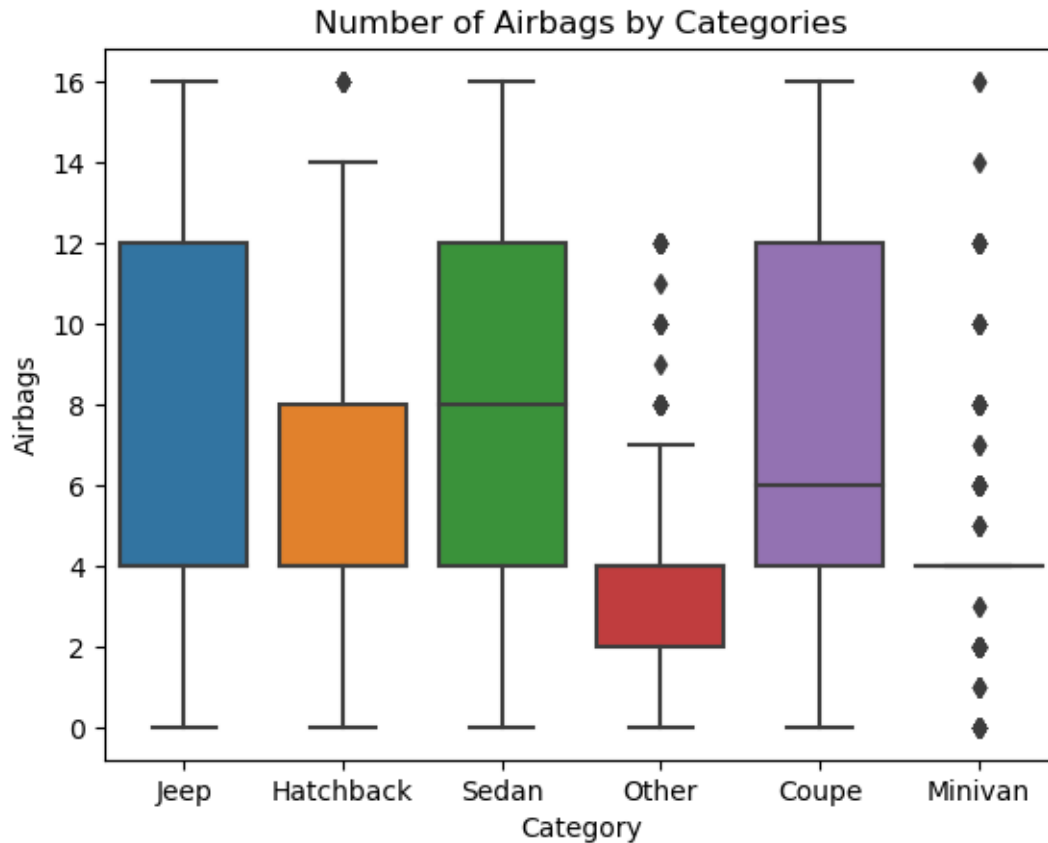


As from the chart above, “Sedan” is the highest sold category in second-hand car sales, followed by “Jeep” and “Hatchback”. The other categories of cars account about a tenth of the total car sales.

```
[69]: #Creating a Boxplot for each category (univariate)
```

```
sns.boxplot(x = "Category" , y = "Airbags" , data = cars_2)
plt.title("Number of Airbags by Categories")

plt.show()
```



From the boxplot above, it can be seen that the average range of airbags count for most categories lies between 4 and 12, with a few outliers extending from 0 to 16, prominently for “Minivan”.

```
[70]: grouped_data_1 = cars.groupby('Year Interval')
```

```
#Calculating the mean value for each group
mean_values_1 = grouped_data_1.mean()
mean_values_1
```

```
[70]:
```

	Price	Levy	Category	Leather interior \
Year Interval				
1930	171.333333	906.299205	3.333333	1.000000
1940	7094.866318	906.299205	7.000000	0.500000
1950	8152.346527	906.299205	4.600000	0.000000



1960	8690.093054	1518.039364	3.800000	0.600000
1970	5094.333333	906.299205	6.555556	0.111111
1980	5164.467742	906.299205	6.903226	0.145161
1990	7113.367273	965.903490	6.345186	0.252385
2000	11898.407163	1052.655317	5.842166	0.600422
2010	15633.366095	839.445091	6.435912	0.825251


Year Interval	Fuel type	Engine volume	Mileage	Cylinders \
1930	5.000000	3.200000	87839.237622	5.333333
1940	5.000000	2.100000	72715.972654	5.000000
1950	5.000000	1.760000	62150.403978	4.000000
1960	4.000000	3.320000	31006.835504	5.600000
1970	4.444444	2.100000	51675.850537	4.777778
1980	3.709677	2.062903	72977.956320	4.290323
1990	3.176930	2.228187	88645.285799	4.615785
2000	3.278802	2.424597	99299.942609	4.847926
2010	3.508377	2.265050	71967.959321	4.466132

Year Interval	Gear box type	Drive wheels	Doors	Wheel	Airbags
1930	0.333333	2.000000	4.000000	1.000000	0.000000
1940	0.500000	2.000000	4.000000	1.000000	0.000000
1950	1.000000	1.200000	3.200000	1.000000	0.200000
1960	0.600000	1.000000	2.800000	1.000000	4.800000
1970	0.888889	1.888889	3.555556	1.000000	2.777778
1980	0.935484	1.274194	3.580645	1.000000	1.112903
1990	0.856028	1.241977	3.666956	0.890720	3.392021
2000	0.754992	0.884601	3.901882	0.829301	6.137865
2010	0.414910	0.886172	3.961764	0.964890	7.078236

[71]: *#Making a Linechart to understand the Price changes over the years (bivariate)*

```
x_values_1 = mean_values_1.index
y_values_1 = mean_values_1['Price']
plt.plot(x_values_1, y_values_1, marker = 'o')
plt.xlabel('Year')
plt.ylabel('Average Price')
plt.title('Average Price over the years')

plt.grid(True)
plt.show()
```



As seen, average car prices have risen over the years with an exception (dip) around the 1970s-1980s, and a steep increase thereafter.

```
[72]: #Making a Treemap to understand the distribution of cars by different
      ↪ manufacturers (multivariate)

fig = px.treemap(data_frame = cars_2, path=["Manufacturer", "Category",
      ↪ "Model"],
                  values = 'Price', title = 'Sales distribution by Manufacturer')
fig.show()
```

As from the treemap above, “Hyundai” seems to be the manufacturer with the most number of cars sold, followed by “Toyota” and “Mercedes”. This exhibits that Hyundai targets a range of customers with varied tastes and preferences.

```
[73]: #Imputing the outliers in 'Price' using the quartile method for cars_2
q1, q3 = np.percentile(cars_2["Price"], [25, 75])
iqr = q3 - q1
lower_bound = q1 - 1.5*iqr
upper_bound = q3 + 1.5*iqr
```

```

#Creating conditions to isolate the outliers
outliers_price_2 = cars_2["Price"][(cars_2["Price"] < lower_bound) |
↳(cars_2["Price"] > upper_bound)]

mean_Price_without_outliers_2 = cars_2["Price"][(cars_2["Price"] >=
↳lower_bound) & (cars_2["Price"] <= upper_bound)].mean()

#Replacing the outliers with the mean_Price_without_outliers_2
cars_2["Price"] = cars_2["Price"].where(~((cars_2["Price"] < lower_bound) |
↳(cars_2["Price"] > upper_bound)), mean_Price_without_outliers_2)

```

```
[74]: grouped_data_2 = cars_2.groupby('Fuel type')
```

```

#Calculating the mean value for each group
mean_values_2 = grouped_data_2.mean()
mean_values_2

```

```
[74]:
```

	ID	Price	Levy	Prod. year	\
Fuel type					
CNG	4.565077e+07	8186.859275	948.270665	1999.880597	
Diesel	4.562996e+07	19380.497469	936.673209	2010.976500	
Hybrid	4.551084e+07	10187.959660	771.696828	2012.179147	
Hydrogen	4.578407e+07	20385.000000	906.299205	2012.000000	
LPG	4.571719e+07	12943.936868	839.868818	2012.004520	
Petrol	4.556194e+07	13564.942696	947.793357	2010.844815	
Plug-in Hybrid	4.544400e+07	22247.181944	687.378412	2013.070588	

	Engine volume	Mileage	Cylinders	Doors	Airbags
Fuel type					
CNG	2.479744	2.352272e+07	4.916844	3.925373	4.703625
Diesel	2.387850	7.167544e+05	4.543250	3.913000	5.436000
Hybrid	2.081153	4.753292e+05	4.273241	3.998587	7.784402
Hydrogen	2.400000	1.168000e+05	6.000000	4.000000	8.000000
LPG	2.236836	2.722707e+05	4.355932	3.992090	4.639548
Petrol	2.357075	1.367855e+06	4.712562	3.897717	6.829327
Plug-in Hybrid	1.657647	1.226019e+05	4.094118	4.000000	9.176471


```
[75]: #Making a Barchart to understand the distribution of cars' Fuel type(s) by
↳their prices (bivariate)
```

```

x_values_2 = mean_values_2.index
y_values_2 = mean_values_2['Price']
sns.barplot(x = x_values_2, y = y_values_2, alpha = 0.9)
plt.xlabel('Fuel Type')
plt.ylabel('Average Price')
plt.title('Average Price for each Fuel type')

```

```
plt.show()
```



As seen in the Bar chart above, cars with the fuel type “Plug-in Hybrid” are the most expensive, on average, followed by “Hydrogen”. Cars with the fuel type “CNG” are the cheapest overall.

3.1 We have performed the necessary data cleaning, wrangling and mining along with EDA, and are prepared to build a machine learning model.

## 4 MILESTONE 3

```
[76]: cars.head()
```

```
[76]:
```

	Price	Levy	Category	Leather interior	Fuel type	Engine volume \
0	13328.0	1399.000000	4	1	2	3.5
1	16621.0	1018.000000	4	0	5	3.0
2	8467.0	906.299205	3	0	5	1.3
3	3607.0	862.000000	4	1	2	2.5
4	11726.0	446.000000	3	1	5	1.3

	Mileage	Cylinders	Gear box type	Drive wheels	Doors	Wheel	\
0	115602.858919	6	0	0	4	1	
1	119328.775637	6	2	0	4	1	
2	124300.807955	4	3	1	4	0	
3	105013.051585	4	0	0	4	1	
4	57116.842759	4	0	1	4	1	

	Airbags	Year Interval
0	12	2000
1	8	2010
2	2	2000
3	0	2010
4	4	2010

```
[77]: cars.dtypes
```

```
[77]: Price          float64
Levy              float64
Category          int32
Leather interior  int32
Fuel type         int32
Engine volume     float64
Mileage           float64
Cylinders         int64
Gear box type     int32
Drive wheels      int32
Doors             int64
Wheel            int32
Airbags          int64
Year Interval     int32
dtype: object
```

**4.0.1** We saw in the Feature Selection process that ‘Levy’ and ‘Leather interior’ variables need to be dropped; We can then work on our chosen Machine Learning models.

```
[78]: cars_3 = cars.drop(['Levy', 'Leather interior'], axis = 1)
```

```
[79]: cars_3.columns
```

```
[79]: Index(['Price', 'Category', 'Fuel type', 'Engine volume', 'Mileage',
        'Cylinders', 'Gear box type', 'Drive wheels', 'Doors', 'Wheel',
        'Airbags', 'Year Interval'],
        dtype='object')
```

```
[80]: predictors_2 = ['Category', 'Fuel type', 'Engine volume', 'Mileage',
        'Cylinders', 'Gear box type', 'Drive wheels', 'Doors', 'Wheel',
```

```

    'Airbags', 'Year Interval']
outcome_2 = 'Price'

```

```

[81]: x = pd.get_dummies(cars_3[predictors_2], drop_first = True)
      y = cars_3[outcome_2]

```

```

[82]: y

```

```

[82]: 0      13328.0
      1      16621.0
      2       8467.0
      3       3607.0
      4      11726.0
      ...
      19232      8467.0
      19233     15681.0
      19234     26108.0
      19235      5331.0
      19236       470.0
      Name: Price, Length: 18922, dtype: float64

```

#### 4.0.2 Training and Test Split

```

[83]: #Splitting the data into training and test set
      train_x, valid_x, train_y, valid_y = train_test_split(x, y, test_size = 0.4,
      ↪random_state = 1)

```

4.1 We will build two predictive models - Linear Regression (with Lasso and Ridge optimization) and Random Forest.

#### 4.2 Linear Regression Algorithm

```

[84]: Cars_LR = LinearRegression()

```

```

[85]: Cars_LR.fit(train_x, train_y)

```

```

[85]: LinearRegression()

```

```

[86]: #Printing the coefficients
      print('intercept:', Cars_LR.intercept_)
      print(pd.DataFrame({'Predictor': x.columns, 'coefficient': Cars_LR.coef_}))

```

```

intercept: -876709.3427277593
      Predictor  coefficient
0      Category -294.347078
1      Fuel type -586.421917
2  Engine volume 1435.465649
3      Mileage   -0.028959

```

```

4         Cylinders  -508.355536
5   Gear box type  2876.507413
6     Drive wheels  1002.106395
7           Doors   701.383172
8           Wheel  5906.281333
9         Airbags  -403.231897
10   Year Interval  442.721882

```

```

[87]: #Making predictions on the test data
y_pred_1 = Cars_LR.predict(valid_x)
y_pred_df = pd.DataFrame({'Predicted': y_pred_1})
y_pred_df

```

```

[87]:          Predicted
0      19435.853672
1      10113.230229
2      20682.767588
3      12073.983795
4      14022.946113
...
7564  10730.177845
7565   4643.017585
7566  11786.637663
7567  20041.002204
7568  13378.501259

[7569 rows x 1 columns]

```

```

[88]: from sklearn.metrics import r2_score

```

```

[89]: #Calculating the R-squared score
r1 = r2_score(valid_y, y_pred_1)
print("R-squared score:", r1)

```

R-squared score: 0.14188155699947191

### 4.3 Random Forest Algorithm

```

[90]: from sklearn.ensemble import RandomForestRegressor

```

```

[91]: #We have taken our hyper-parameter (number of decision trees) as 100 to try to
      ↪improve the model performance
rf_model = RandomForestRegressor(n_estimators = 100, random_state = 1)

```

```

[92]: rf_model.fit(train_x, train_y)

```

```

[92]: RandomForestRegressor(random_state=1)

```

```
[93]: #Making predictions on the test data
y_pred_2 = rf_model.predict(valid_x)
y_pred_2df = pd.DataFrame({'Predicted': y_pred_2})
y_pred_2df
```

```
[93]:      Predicted
0      30851.077916
1       545.850000
2      8900.427333
3     10847.730000
4     16310.527560
...
7564  15170.390000
7565   7658.964167
7566  15508.500000
7567  16243.353571
7568  17898.200080

[7569 rows x 1 columns]
```

```
[94]: #Calculating the R-squared score
r2 = r2_score(valid_y, y_pred_2)
print("R-squared score:", r2)
```

R-squared score: 0.5357291795068264

#### 4.4 Lasso (Linear Regression) Algorithm

```
[95]: from sklearn.linear_model import Lasso
```

```
[96]: #We have taken our hyper-parameter (regularization strength) as 0.1 to try to
      ↪improve the model performance by penalizing features
lasso_model = Lasso(alpha = 0.1)
```

```
[97]: lasso_model.fit(train_x, train_y)
```

```
[97]: Lasso(alpha=0.1)
```

```
[98]: #Making predictions on the test data
y_pred_3 = lasso_model.predict(valid_x)
y_pred_3df = pd.DataFrame({'Predicted': y_pred_3})
y_pred_3df
```

```
[98]:      Predicted
0     19435.512099
1     10113.197278
2     20683.542968
3     12073.876024
```



```

4      14023.121445
...
7564   10730.559644
7565    4644.490640
7566   11786.700788
7567   20040.667372
7568   13378.220397

```

[7569 rows x 1 columns]

```

[99]: #Calculating the R-squared score
      r3 = r2_score(valid_y, y_pred_3)
      print("R-squared score:", r3)

```

R-squared score: 0.14188334035891337

## 4.5 Ridge (Linear Regression) Algorithm

```

[100]: from sklearn.linear_model import Ridge

```

```

[101]: #We have taken our hyper-parameter (regularization strength) as 1 to try to
      ↪improve the model performance by penalizing features
      ridge_model = Ridge(alpha = 1)

```

```

[102]: ridge_model.fit(train_x, train_y)

```

```

[102]: Ridge(alpha=1)

```

```

[103]: #Making predictions on the test data
      y_pred_4 = ridge_model.predict(valid_x)
      y_pred_4df = pd.DataFrame({'Predicted': y_pred_4})
      y_pred_4df

```

```

[103]:      Predicted
0      19435.306770
1      10113.416212
2      20688.334992
3      12073.754776
4      14022.485274
...
7564   10730.283237
7565    4650.003525
7566   11786.452580
7567   20039.433361
7568   13378.389630

```

[7569 rows x 1 columns]

```
[104]: #Calculating the R-squared score
r4 = r2_score(valid_y, y_pred_4)
print("R-squared score:", r4)
```

R-squared score: 0.14189020110392536

R-squared ( $R^2$ ) is a statistical metric used to evaluate the goodness of fit of a regression model. It provides a measure of how well the independent variables (features) explain the variability of the dependent variable (target) in the regression model. R-squared is also known as the coefficient of determination.

Mean Error (ME): The average difference between the predicted and actual values, indicating the overall bias of the model's predictions.

Root Mean Squared Error (RMSE): The square root of the average of the squared differences between predicted and actual values, representing the model's overall accuracy with a focus on larger errors.

Mean Absolute Error (MAE): The average of the absolute differences between predicted and actual values, providing a measure of the model's overall accuracy without considering the direction of errors.

Mean Percentage Error (MPE): The average percentage difference between predicted and actual values, indicating the model's overall bias in percentage terms.

Mean Absolute Percentage Error (MAPE): The average percentage difference between predicted and actual values, providing a relative measure of the model's accuracy.

#### 4.5.1 The best metrics to evaluate model performance are R-squared, RMSE and MAE.

### 4.6 Regression Summary for Training Data

```
[105]: #Simple Linear Regression
regressionSummary(train_y, Cars_LR.predict(train_x))
```

Regression statistics

```
Mean Error (ME) : 0.0000
Root Mean Squared Error (RMSE) : 10229.2364
Mean Absolute Error (MAE) : 7963.2868
Mean Percentage Error (MPE) : -1372.4081
Mean Absolute Percentage Error (MAPE) : 1407.1201
```

```
[106]: #Random Forest
regressionSummary(train_y, rf_model.predict(train_x))
```

Regression statistics

```
Mean Error (ME) : 17.4339
```

```
Root Mean Squared Error (RMSE) : 3037.3739
Mean Absolute Error (MAE) : 1877.3106
Mean Percentage Error (MPE) : -339.6282
Mean Absolute Percentage Error (MAPE) : 348.3782
```

```
[107]: #Lasso Regression
regressionSummary(train_y, lasso_model.predict(train_x))
```

Regression statistics

```
Mean Error (ME) : -0.0000
Root Mean Squared Error (RMSE) : 10229.2364
Mean Absolute Error (MAE) : 7963.2566
Mean Percentage Error (MPE) : -1372.4356
Mean Absolute Percentage Error (MAPE) : 1407.1450
```

```
[108]: #Ridge Regression
regressionSummary(train_y, ridge_model.predict(train_x))
```

Regression statistics

```
Mean Error (ME) : 0.0000
Root Mean Squared Error (RMSE) : 10229.2366
Mean Absolute Error (MAE) : 7963.1661
Mean Percentage Error (MPE) : -1372.5140
Mean Absolute Percentage Error (MAPE) : 1407.2194
```

## 4.7 Regression Summary for Validation Data

```
[109]: #Simple Linear Regression
regressionSummary(valid_y, Cars_LR.predict(valid_x))
```

Regression statistics

```
Mean Error (ME) : -254.5642
Root Mean Squared Error (RMSE) : 10169.6145
Mean Absolute Error (MAE) : 7934.0439
Mean Percentage Error (MPE) : -1580.6005
Mean Absolute Percentage Error (MAPE) : 1616.1786
```

```
[110]: #Random Forest
regressionSummary(valid_y, rf_model.predict(valid_x))
```

Regression statistics

```
Mean Error (ME) : -188.1848
```

```
Root Mean Squared Error (RMSE) : 7480.2579
Mean Absolute Error (MAE) : 4868.6604
Mean Percentage Error (MPE) : -978.1230
Mean Absolute Percentage Error (MAPE) : 999.6086
```

```
[111]: #Lasso Regression
regressionSummary(valid_y, lasso_model.predict(valid_x))
```

Regression statistics

```
Mean Error (ME) : -254.5767
Root Mean Squared Error (RMSE) : 10169.6040
Mean Absolute Error (MAE) : 7934.0087
Mean Percentage Error (MPE) : -1580.6291
Mean Absolute Percentage Error (MAPE) : 1616.2065
```

```
[112]: #Ridge Regression
regressionSummary(valid_y, ridge_model.predict(valid_x))
```


Regression statistics


```
Mean Error (ME) : -254.5805
Root Mean Squared Error (RMSE) : 10169.5633
Mean Absolute Error (MAE) : 7933.9005
Mean Percentage Error (MPE) : -1580.6605
Mean Absolute Percentage Error (MAPE) : 1616.2362
```

```
[113]: #Plotting the graphs for the two main models
plt.scatter(x = valid_y, y = y_pred_1, color = 'blue', label = 'Actual vs_
↳ Predicted - Test Set')
plt.plot([min(y), max(y)], [min(y), max(y)], color = 'red', linestyle='--',
↳ label = 'Perfect Fit')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Linear Regression')
plt.legend()
plt.show()

plt.scatter(x = valid_y, y = y_pred_2, color = 'blue', label = 'Actual vs_
↳ Predicted - Test Set')
plt.plot([min(y), max(y)], [min(y), max(y)], color = 'red', linestyle='--',
↳ label = 'Perfect Fit')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Random Forest Regressor')
plt.legend()
```

```
plt.show()
```





- 4.8 From the Regression summary metrics above, we can see that the Linear Regression models have higher error as compared to the Random Forest Regressor. Therefore, it is in the best interest of ABC Motors to use a Random Forest algorithm in order to predict car prices and improve their profitability.